
**Module 07 Assignment 01: Programming Assignment 2**

**Name: Advaithbarath Raghuraman Bhuvaneswari**

**UCID: ar2728**

**GitHub:** rbadvaith/Wine_Cluster

**Docker:** advaith123/bucketcluster Tags | Docker Hub

## Short Summary

This assignment involves building a machine learning application for predicting wine quality using Apache Spark on the Amazon AWS cloud platform. The goal is to train a model in parallel across multiple EC2 instances and deploy it effectively. Using provided datasets, the application will first train a model on Spark's MLlib in parallel on four EC2 instances, validate its performance using a second dataset, and optimize parameters for the best accuracy. The trained model will then be used on a single EC2 instance for predictions, with its performance measured by the F1 score. Additionally, the prediction application must be containerized using Docker to enable efficient deployment in diverse environments. The implementation is to be completed in Java on an Ubuntu Linux setup, integrating skills in parallel computing, ML model development, and containerization.

## Initial steps for AWS Credentials setup

The credentials created here will be used by our application for connecting to Elasticmapreduce service and S3.

Via Aws Learner(If you are a new user)

1. Create an AWS Account with NJIT email.
2. Go to *Academy Learner lab > Modules>Launch Learner Academy Lab*
3. Learner lab is being displayed here

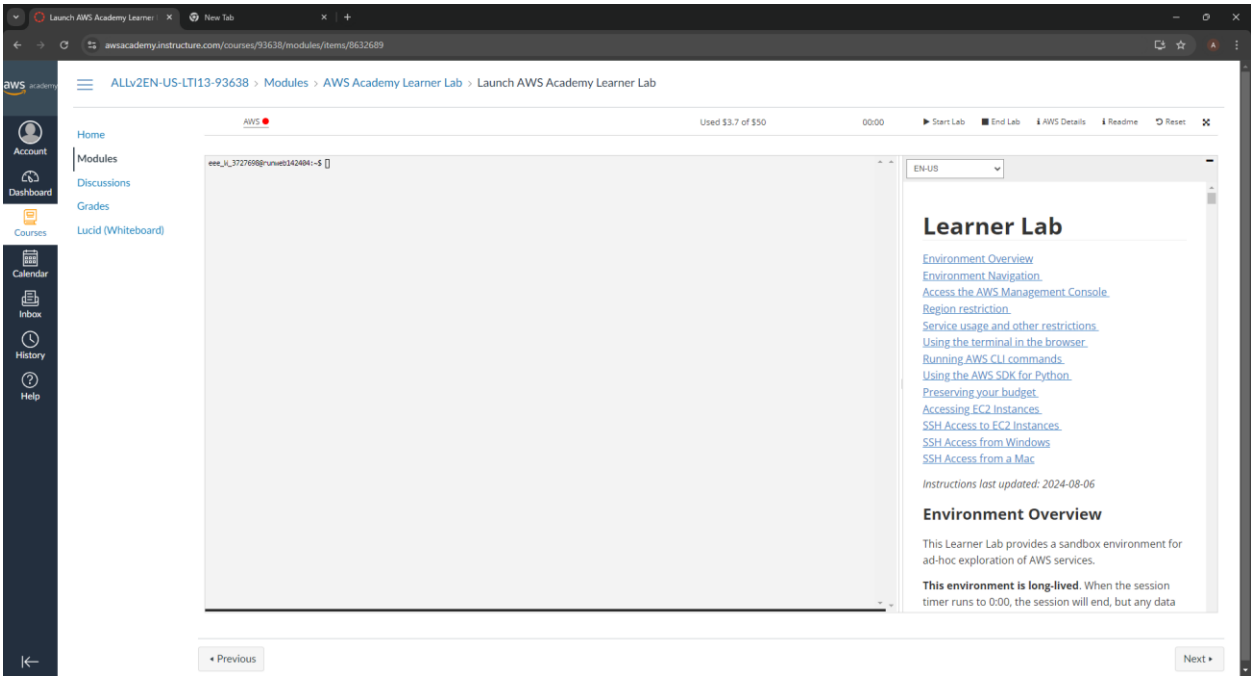


Fig 1. Learner Lab via AWS Academy

Via Vocareum(If you are an existing user | Use Microsoft Edge)

1. If you are an existing user, log in via NJIT email.
2. Select role as student, the search result will display a lab link. Click that to open .
3. Learner lab is being displayed here

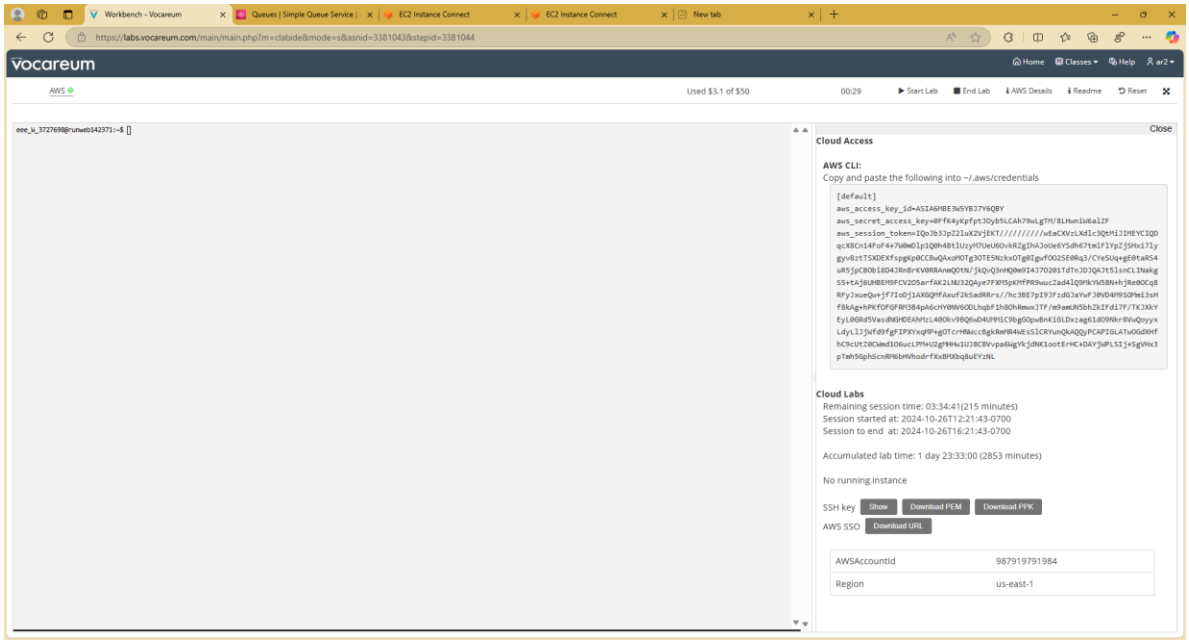


Fig 2. Learner Lab via AWS Academy

## Launching AWS Management Console

1. Click Start Lab on the right-hand corner
2. Wait for one minute to see the AWS symbol turned from red to green
3. Click the AWS symbol once it turned green
4. AWS Management Console is being displayed. Please note that the session will last for 4 hrs. only

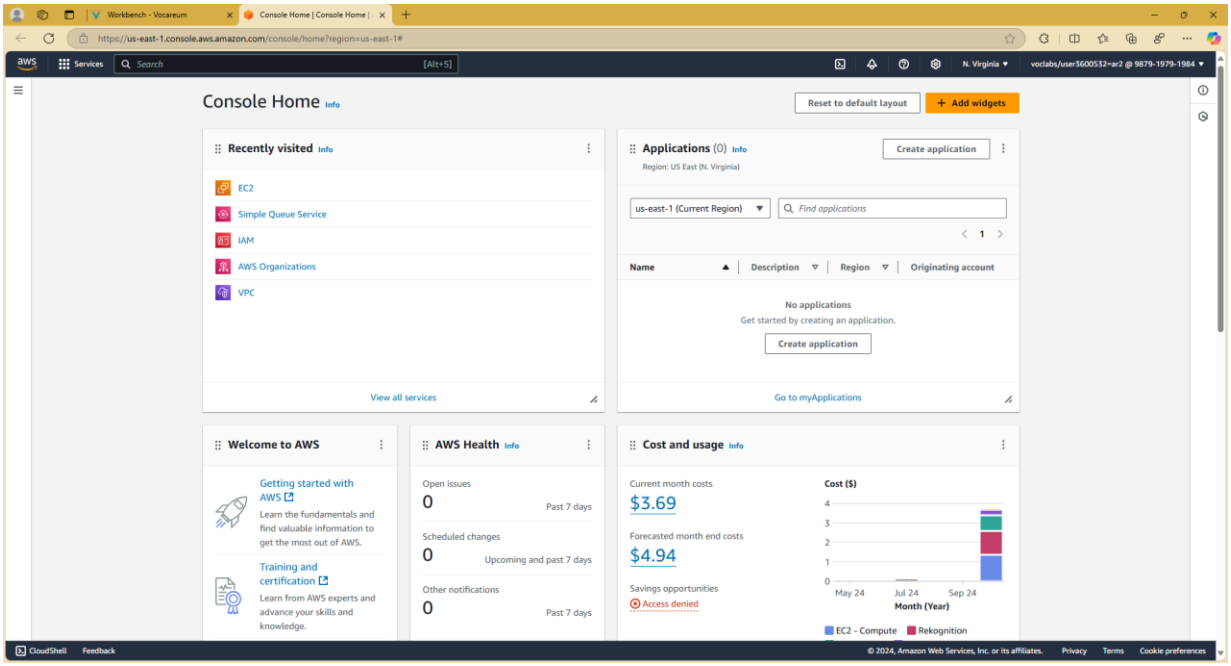


Fig 3. AWS Management Console

## Creating EMR Cluster

EMR Cluster 1 core and 4 slaves will be created for the Elasticmapreduce service. These service will only run until the session end. After session finished, you cant able to access the cluster. But don’t worry, you can clone the old terminated cluster and re-run it again.

1. In the search bar, type "*EMR*" and then select the first result.
2. Select "*Create cluster*" from the navigation menu on the right. If you want to clone the existing cluster, select your preferred cluster and select “*clone*”

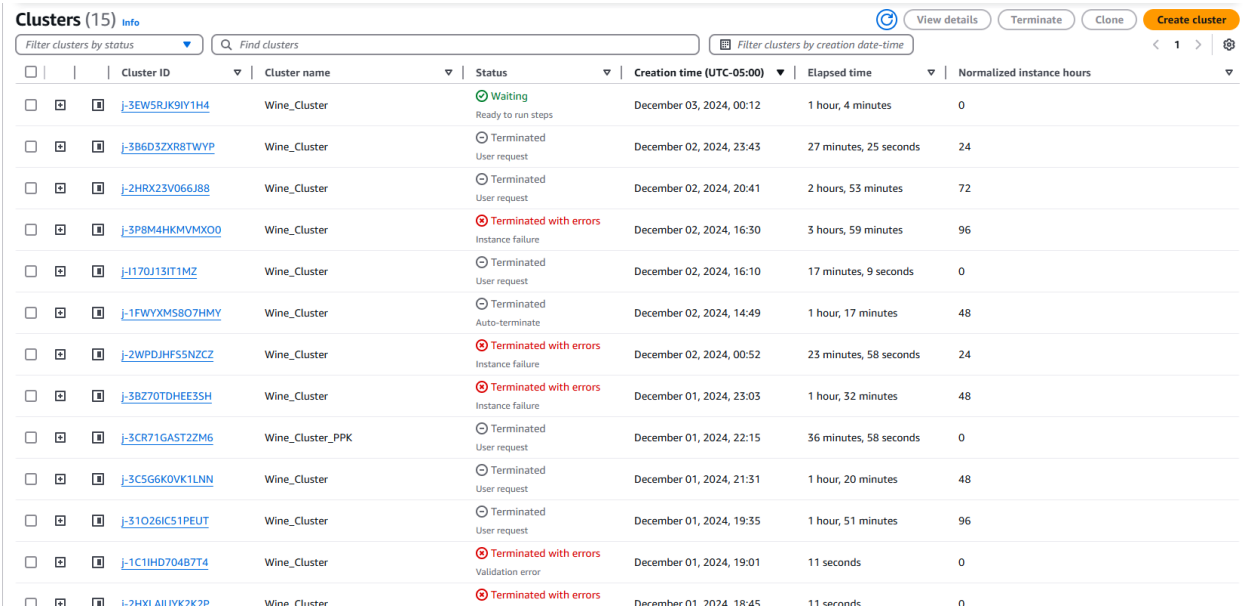


Fig 4. Create Cluster

3. In *Name*: Type Wine_Cluster
4. Ensure that the application bundle should be the same as in Fig.5

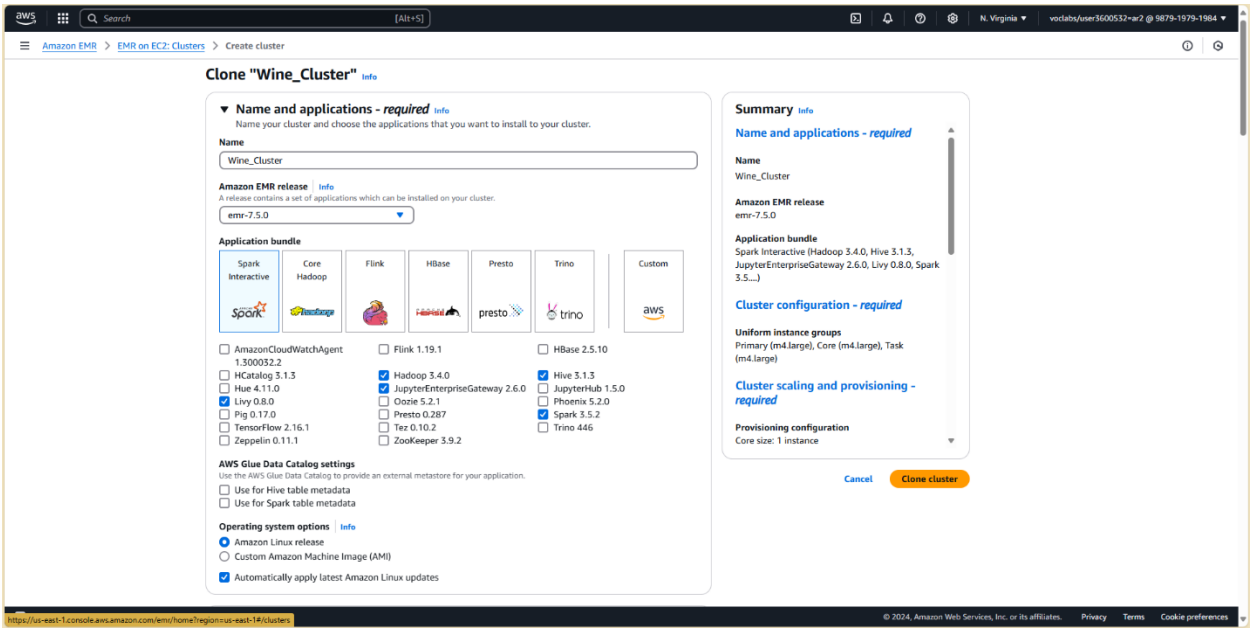

Fig 5. Clone or Create Cluster

5. Select instance type as m4large for all

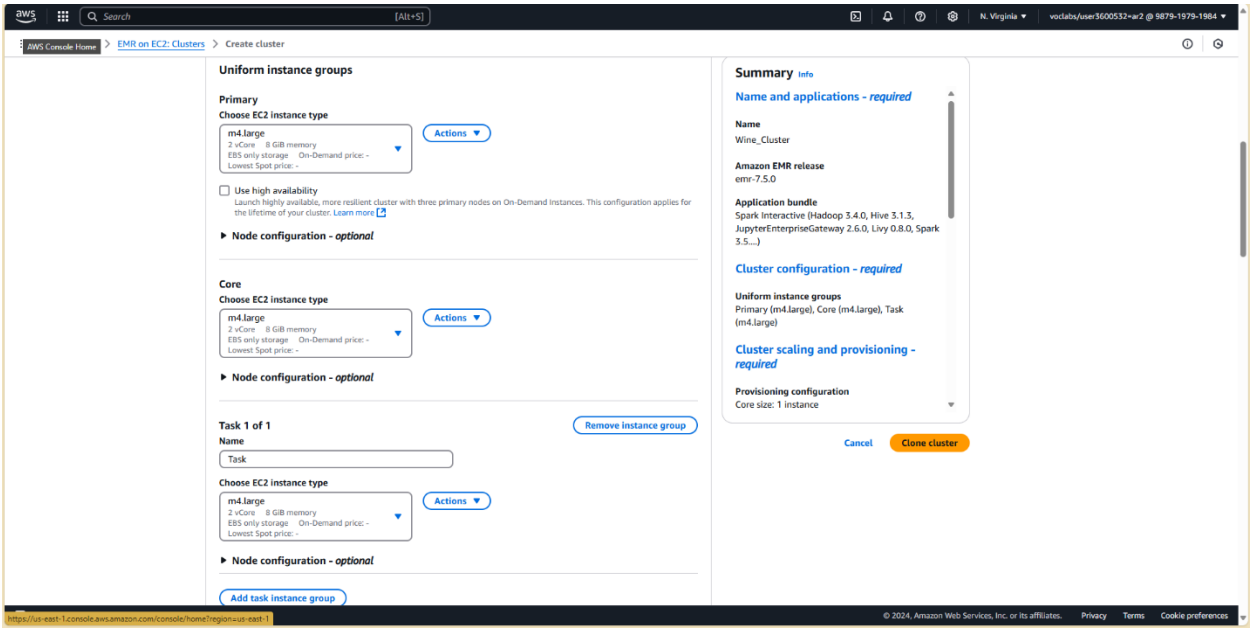

Fig 6. Instance Type

6. In Cluster Termination, set termination to manual from automatic

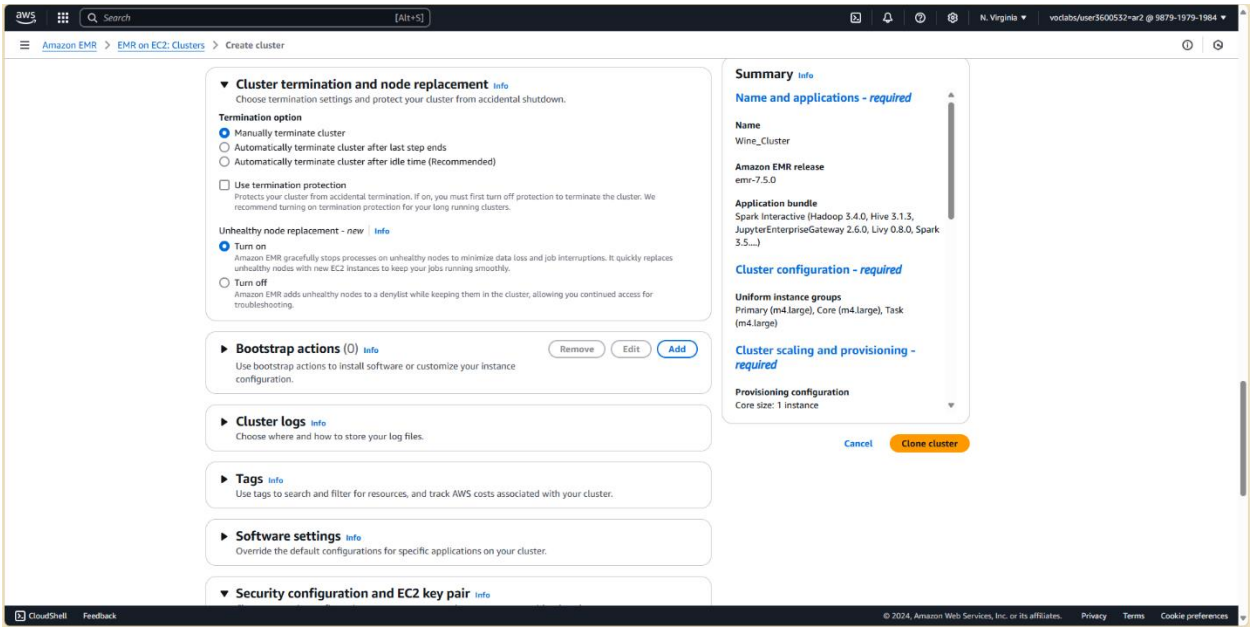


Fig 7. Cluster Termination

7. Set EC2 Private key and name it as cluster.ppk

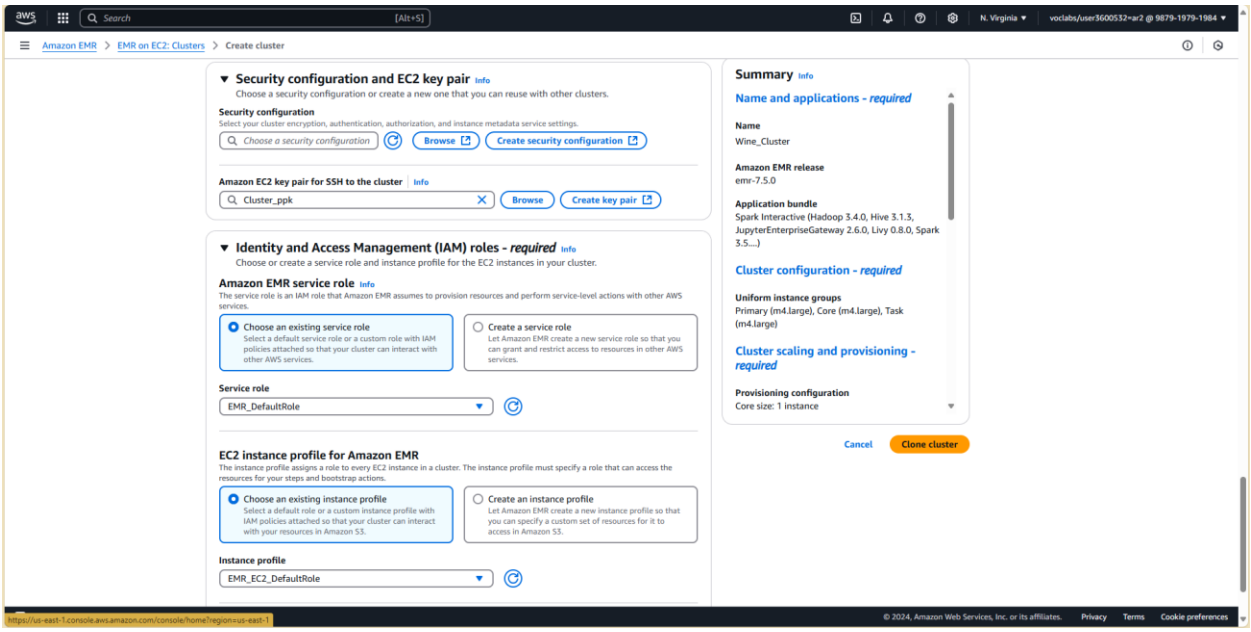


Fig 8. EC2 Key Pair

8. Select Create Cluster or Clone Cluster and wait for 5 minutes to load

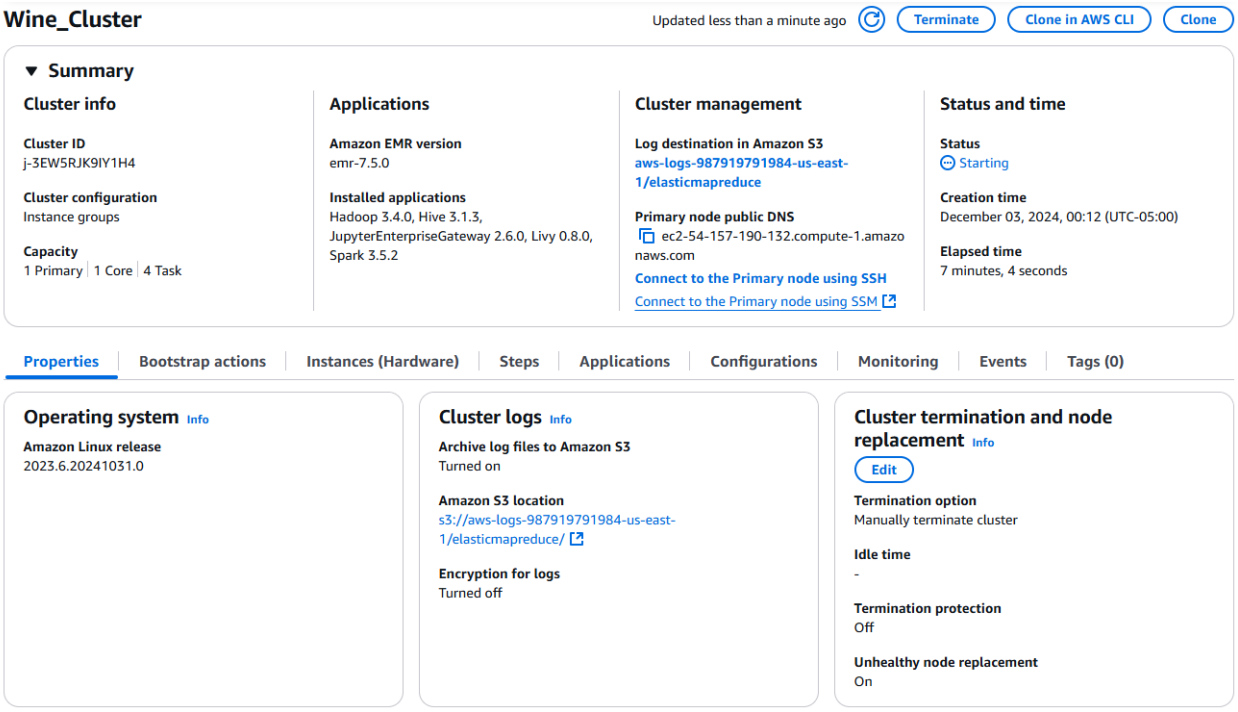


Fig 9. Wine_Cluster

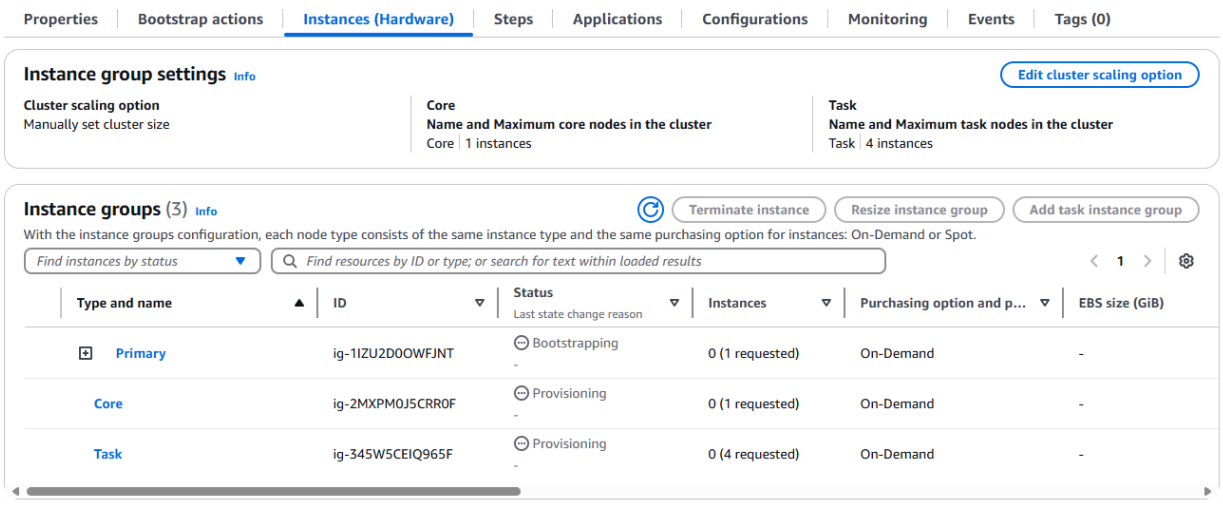


Fig 10. Number of Instances

## Setting SSH for master(core) EC2 in the cluster[Only required to do this for first time]

We need to provide ssh access for the master instance which is not provided when initially created.

1. Search EC2 in the search bar and click enter
2. In the Dashboard select security group
3. Over there Select the group id that have group name elasticmapreduce master in it

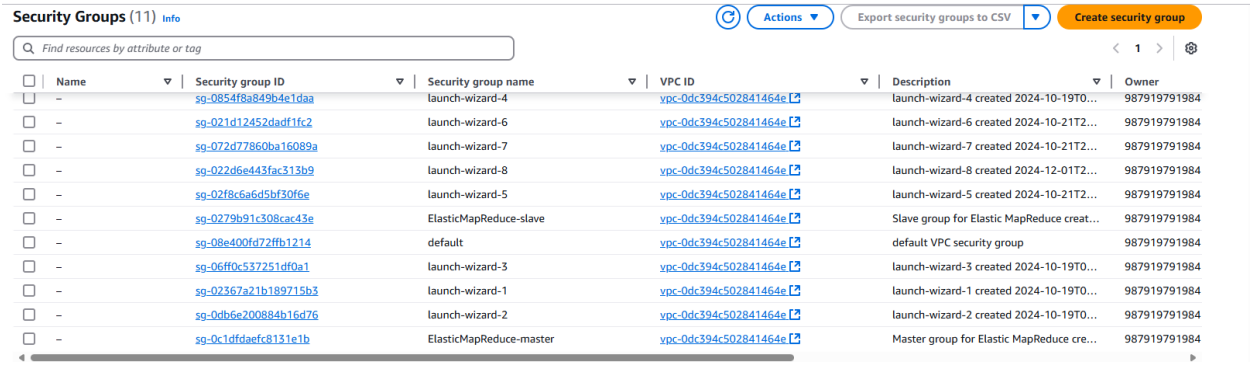


Fig 11. Security Group Name

4. Select edit inbound rules in the middle right edge

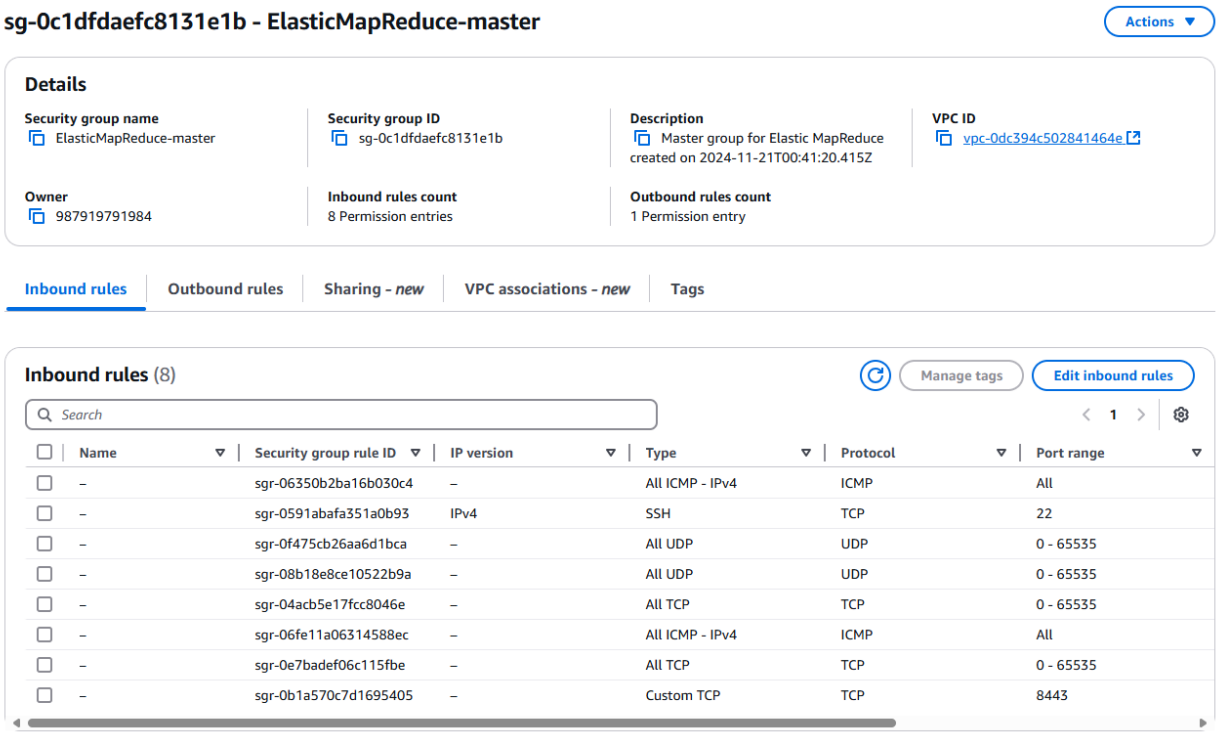


Fig 12. ElasticMapReduce-master

5. Verify if ssh is present there else, create a new one

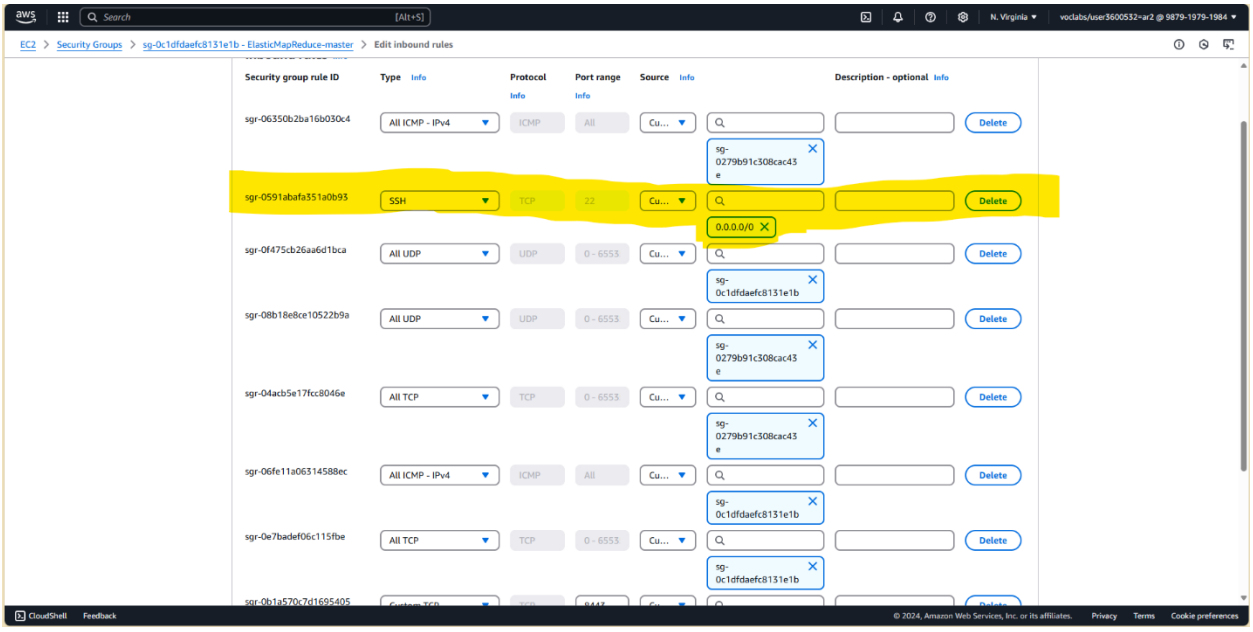


Fig 13. SSH in Inbound rules

6. Now you could able to access the master instance, now to EMR menu, select connect via ssm, go to EC2 connect tab and click connect

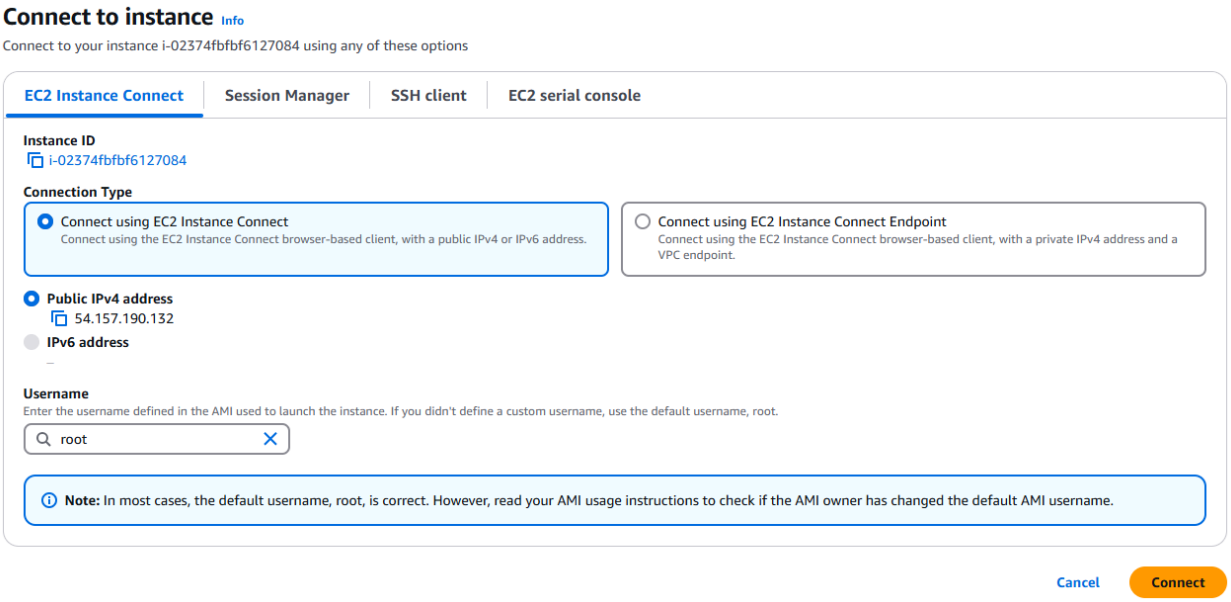


Fig 14. Connect to EC2

## S3 Bucket

We are going to store the code and csv file in S3 bucket to sync them and use it in EMR Cluster

1. In the search bar, type S3 and select create bucket.
2. Give name *buckercluster* to it
3. After creating it, *upload training.py, TrainingDataset.csv, prediction.py* and *ValidationDataset.csv* to this bucket
4. Check whether is these file present as given in below figure. Pls note that *trainedmodel* folder is also present in the bucket which is created only after training.py have finished execution in the EMR Cluster

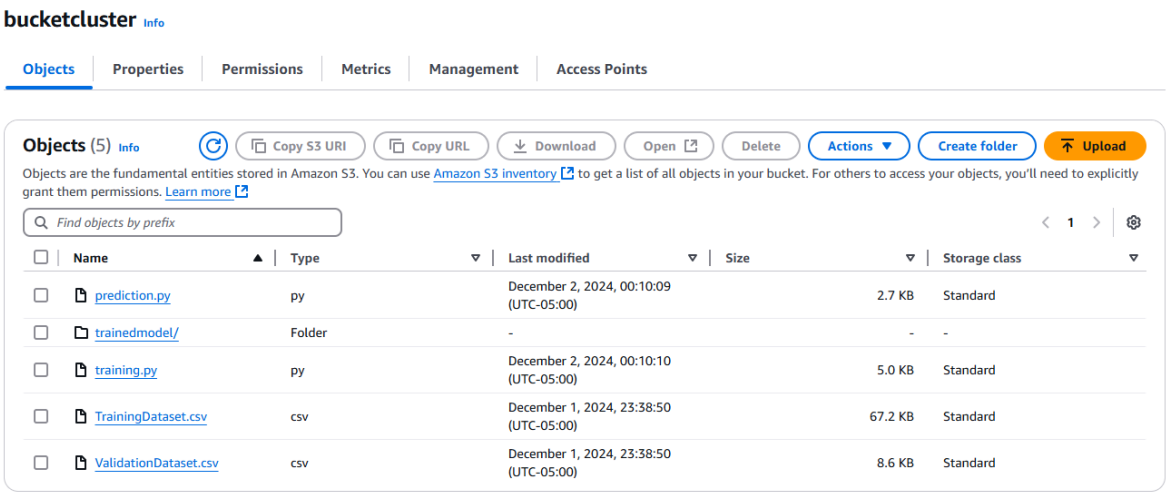


Fig 15. Bucketcluster

## Execution without Docker

We are going to execute both training.py and prediction.py without docker

1. Go to EC2 Instance
2. Install numpy via cmd "***pip install numpy***"

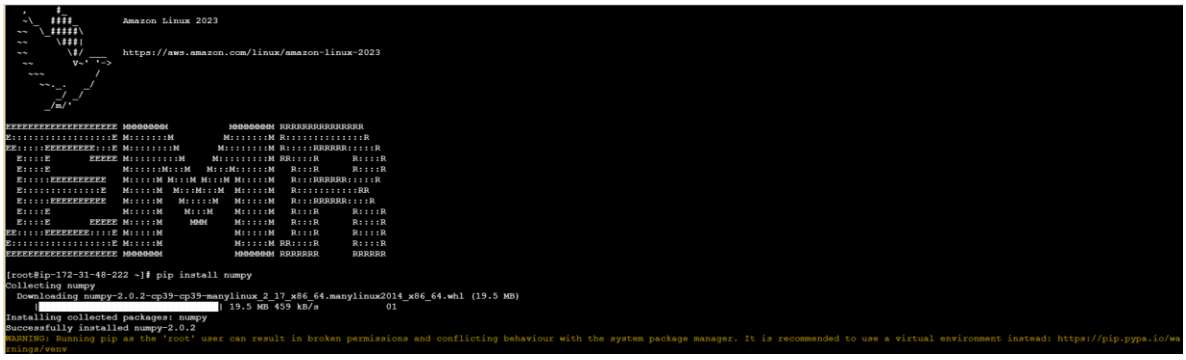


Fig 16. Install numpy

3. Now Run training.py via command "***spark-submit s3://bucketcluster/training.py***" and wait until execution get finished

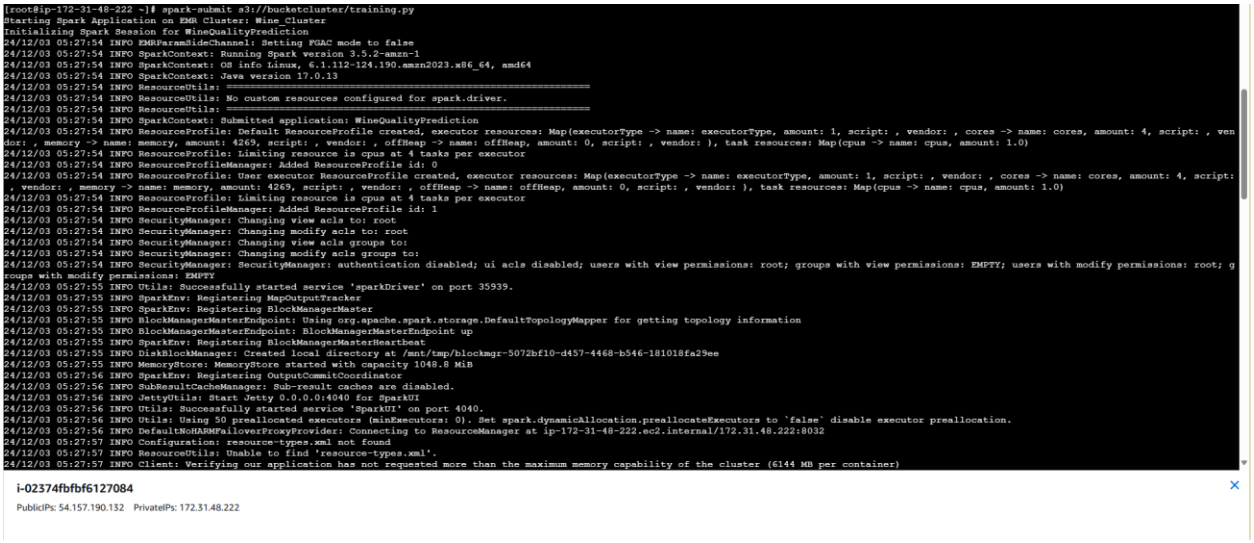


Fig 17. Running training.py


Fig 18. Completed executing training.py

4. Now Run prediction.py via command "***spark-submit s3://bucketcluster/prediction.py s3://bucketcluster/ValidationDataset.csv***" and wait until execution get finished


Fig 19. Running prediction.py


Fig 20. Completed executing prediction.py

5. We could able to conclude that after running prediction.py, we could able to find accuracy score is 96% and F1 score is 0.954791

## Execution with Docker

We are going to execute prediction.py with docker

1. Go to docker desktop, login there
2. Build, tag and push the compiled prediction.py images into docker so that it could be accessed by EC2. We need to do this step only if there any changes in the code or if you are pushing it for the first time.

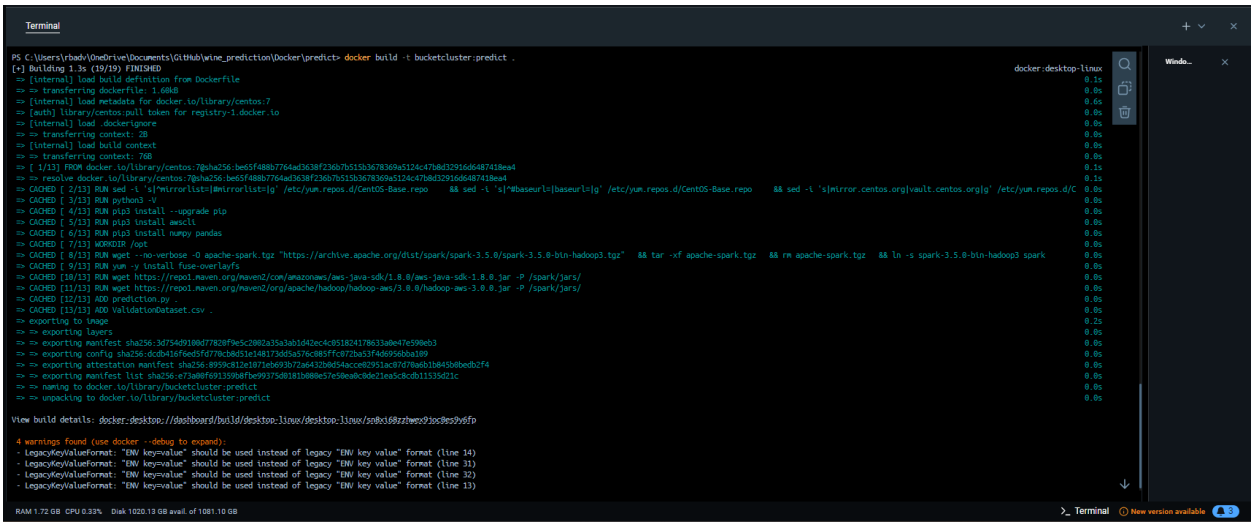


Fig 21. Building Image of prediction.py in docker

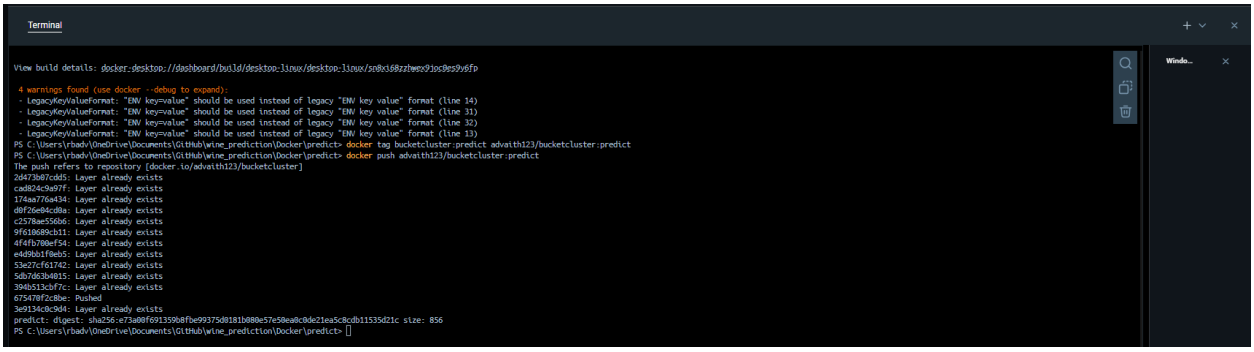


Fig 22. Tagging and Pushing Image of prediction.py in docker

3. Now go to EC2 instance and login into docker there. Also type cmd "***sudo systemctl start docker*** " and " ***sudo systemctl enable docker***" as well

```
[root@ip-172-31-48-222 ~]# docker login
Authenticating with existing credentials...
WARNING! Your password will be stored unencrypted in /root/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
[root@ip-172-31-48-222 ~]# sudo systemctl start docker
[root@ip-172-31-48-222 ~]# sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
```

Fig 23. Starting Docker in EC2

4. Pull prediction.py image from docker using command "***sudo docker pull advaith123/bucketcluster:predict*** ". The below Figure states that the images have already been pulled

```
[root@ip-172-31-48-222 ~]# sudo docker pull advaith123/bucketcluster:predict
predict: Pulling from advaith123/bucketcluster
Digest: sha256:b65a3609d143687b7226dafb8f1e117c619ef3a51eb8d46bcf0275977275f0f5
Status: Image is up to date for advaith123/bucketcluster:predict
docker.io/advaith123/bucketcluster:predict
```

Fig 24. Pulling images to EC2

5. Use command "***docker images***" to check whether is image uploaded or not

```
[root@ip-172-31-48-222 ~]# docker images
REPOSITORY                  TAG       IMAGE ID       CREATED        SIZE
advaith123/bucketcluster    predict   dcdb416f6ed5   2 hours ago    1.56GB
```

Fig 25. Checking whether if images exists in EC2

6. Now Run prediction.py via command "***sudo docker run -v /home/ec2-user/:/job advaith123/bucketcluster:predict***" and wait until execution get finished

```
[root@ip-172-31-48-222 ~]# sudo docker run -v /home/ec2-user/:/job advaith123/bucketcluster:predict
:: loading settings :: url = jar:file:/opt/spark-3.5.0-bin-hadoop3/jars/ivy-2.5.1.jar!/org/apache/ivy/core/settings/ivysettings.xml
Ivy Default Cache set to: /root/.ivy2/cache
The jars for the packages stored in: /root/.ivy2/jars
org.apache.hadoop#hadoop-aws added as a dependency
com.amazonaws#aws-java-sdk added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-852f78b7-2caa-417d-8f2e-ee38d915ee8f;1.0
        confs: [default]
        found org.apache.hadoop#hadoop-aws;2.7.3 in central
        found org.apache.hadoop#hadoop-common;2.7.3 in central
        found org.apache.hadoop#hadoop-annotations;2.7.3 in central
        found com.google.guava#guava;11.0.2 in central
        found com.google.code.findbugs#jsr305;3.0.0 in central
        found commons-cli#commons-cli;1.2 in central
        found org.apache.commons#commons-math3;3.1.1 in central
        found xmlenc#xmlenc;0.52 in central
        found commons-httpclient#commons-httpclient;3.1 in central
        found commons-logging#commons-logging;1.1.3 in central
        found commons-codec#commons-codec;1.4 in central
        found commons-io#commons-io;2.4 in central
        found commons-net#commons-net;3.1 in central
        found commons-collections#commons-collections;3.2.2 in central
        found javax.servlet#servlet-api;2.5 in central
        found org.mortbay.jetty#jetty;6.1.26 in central
        found org.mortbay.jetty#jetty-util;6.1.26 in central
        found com.sun.jersey#jersey-core;1.9 in central
        found com.sun.jersey#jersey-json;1.9 in central
        found org.codehaus.jettison#jettison;1.1 in central
        found com.sun.xml.bind#jaxb-impl;2.2.3-1 in central
        found javax.xml.bind#jaxb-api;2.2.2 in central
        found javax.xml.stream#stax-api;1.0-2 in central
        found javax.activation#activation;1.1 in central
        found org.codehaus.jackson#jackson-core-asl;1.9.13 in central
        found org.codehaus.jackson#jackson-mapper-asl;1.9.13 in central
        found org.codehaus.jackson#jackson-jaxrs;1.9.13 in central
        found org.codehaus.jackson#jackson-xc;1.9.13 in central
        found com.sun.jersey#jersey-server;1.9 in central
        found asm#asm;3.2 in central
        found log4j#log4j;1.2.17 in central
        found net.java.dev.jets3t#jets3t;0.9.0 in central
        found org.apache.httpcomponents#httpclient;4.2.5 in central
```

Fig 26. Running Docker Prediction Image in EC2


Fig 27. Completed executing Docker Prediction Image in EC2

7. We could able to conclude that after running prediction image from docker, we could able to find accuracy score is 96% and F1 score is 0.954791

## Docker Image

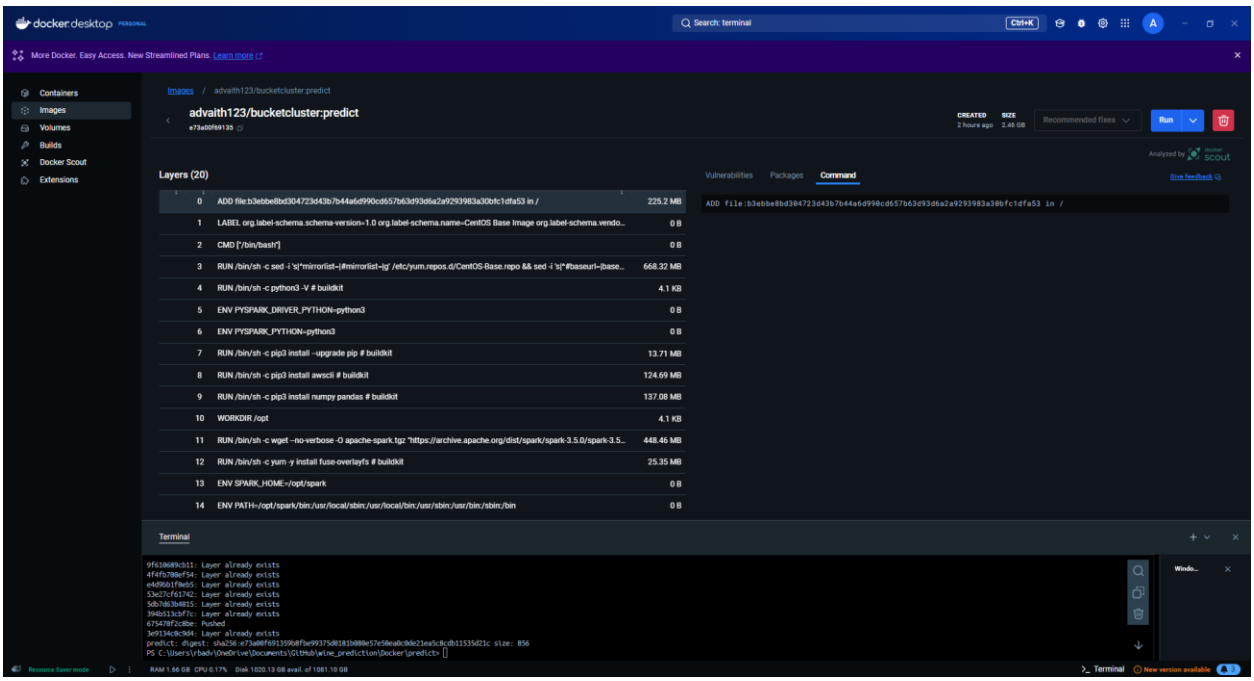


Fig 28. Docker Image